

Lecture 10

Threads and Concurrency

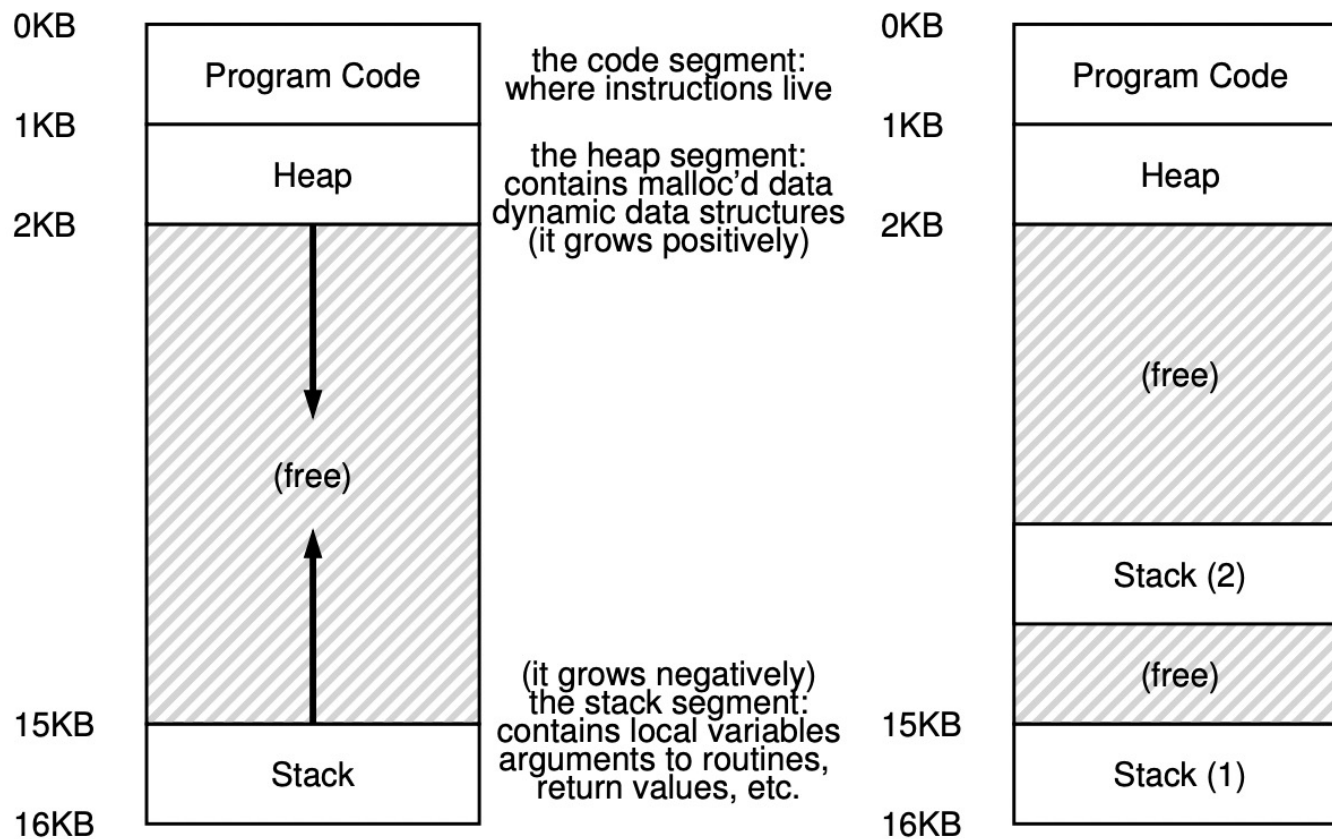
Prof. Yinqian Zhang

Spring 2022

What is a Thread?

- Thread is an **abstraction** of the execution of a program
 - A single-threaded program has one point of execution
 - A multi-threaded program has more than one points of execution
- Each thread has its own **private** execution state
 - Program counter and a private set of registers
 - A private stack for thread-local storage
 - CPU switching from one thread to another requires context switch
- Threads in the same process **share** computing resources
 - Address space, files, signals, etc.

Single-Threaded and Multi-Threaded



Why Use Thread?


- Increase parallelism
 - One thread per CPU makes better use of multiple CPUs to improve efficiency
- Avoid blocking program progress due to slow I/O
 - Threading enables overlap of I/O with other activities within a single program
 - e.g., many modern server-based applications (web servers, database management systems, and the like) make use of threads
- And allow resource sharing !!!

Thread Implementation


- User-level thread
 - Thread management (e.g., creating, scheduling, termination) done by user-level threads library
 - OS does not know about user-level thread
- Kernel-level thread
 - Thread management done by kernel
 - OS is aware of each kernel-level thread

Thread Models


- One-to-one mapping
 - One user-level thread to one kernel-level thread
- Many-to-one mapping
 - Many user-level thread to one kernel-level thread
- Many-to-many mapping
 - Many user-level thread to many kernel-level thread



One-to-One



Many-to-One



Many-to-Many

Pros and Cons

- Many-to-one mapping
 - Pros: context switch between threads is cheap
 - Cons: When one thread blocks on I/O, all threads block
- One-to-one mapping
 - Pros: Every thread can run or block independently
 - Cons: Need to make a crossing into kernel mode to schedule
- Many-to-many mapping
 - Many user-level threads multiplexed on less or equal number of kernel-level threads
 - Pros: best of the two worlds, more flexible
 - Cons: difficult to implement

POSIX Thread

- A POSIX standard API for thread creation and synchronization
 - **Portable Operating System Interface (POSIX)** is a family of standards specified by the IEEE Computer Society for maintaining compatibility between operating systems.
 - POSIX APIs are defined for software compatibility on variants of Unix and other operating systems.
 - POSIX APIs are common in UNIX-like operating systems (e.g., Solaris, Linux, Mac OS X)
- The API specifies behavior of the thread library, but the implementation is up to development of the library

Process-wide Attributes

- POSIX thread may be implemented either as user-level or kernel-level
- POSIX thread requires threads of the same process share
 - process ID
 - parent process ID
 - process group ID and session ID
 - user and group IDs
 - open file descriptors
 - current directory
 - resource limits
 - measurements of the consumption of CPU time and resources

Pthread Programming

```
/* thread.c */
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
int sum = 0; //shared by all threads
void *runner(void *param) {
    int i, upper = atoi(param);
    for (i=1;i<=upper;i++)
        sum += i;
    printf("PID: %d, PPID: %d\n", (int)getpid(), (int)getppid());
    pthread_exit(0);
}
int main() {
    pthread_t tid; // thread identifier
    pthread_attr_t attr; //set of thread attributes
    pthread_attr_init(&attr); //get the default attributes
    printf("PID: %d, PPID: %d\n", (int)getpid(), (int)getppid());
    /* create the thread */
    pthread_create(&tid, &attr, runner, "10");
    /* wait for the thread to exit */
    pthread_join(tid, NULL);
    printf("sum = %d\n", sum);
}

/* gcc -o test thread.c -lpthread */
```

Thank you!

